

project_eda

```
library(dplyr)

## Warning: package 'dplyr' was built under R version 3.3.2
##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##     filter, lag
##
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union

# lab = read.csv("data/health/labs.csv")
demographics = read.csv("data/health/demographic.csv")
ques = read.csv("data/health/questionnaire.csv")
exam = read.csv("data/health/examination.csv")

# join tables by participant ID
df = inner_join(demographics, ques, by="SEQN")
df = inner_join(df, exam, by="SEQN")
dim(df)

## [1] 9813 1222

# subset df_adult to consider only the variables we're interested in
response = "BMXBMI"
predictors = c("RIAGENDR", "RIDAGEYR", "RIDRETH3", "DMDEDUC2", "DMDMARTL", "DMDFMSIZ", "INDFMIN2",
               "ALQ101", "ALQ120Q",
               "CBD070", "CBD120", "CBD130",
               "DBD895", "DBD900", "DBQ197",
               "DPQ020", "DPQ030",
               "PAQ710",
               "SLD010H", "SMQ040")

columns = c(predictors, response)

df = df[names(df) %in% columns]

# rename columns to more intuitive names
df = rename(df, gender = RIAGENDR, age = RIDAGEYR, race = RIDRETH3, edu = DMDEDUC2,
            marriage = DMDMARTL, famsize = DMDFMSIZ, famincome = INDFMIN2,
            alcohol12yr = ALQ101, alcoholfrq = ALQ120Q, grocery = CBD070, eatout = CBD120,
            delivery = CBD130, milk = DBQ197, meals_nothome = DBD895, meals_fastfood = DBD900,
            depressed = DPQ020, sleep_trouble = DPQ030, tv_hrs = PAQ710,
            sleep_hr = SLD010H, smoke = SMQ040, bmi = BMXBMI)

# subset the df to consider only adults aged 18 or above
df_adult = df[df$age > 20,]

df_adult = df_adult[which(df_adult$grocery!=777777 & df_adult$grocery!=999999),]
```

```

df_adult = df_adult[which(df_adult$eatout!=77777 & df_adult$eatout!=999999),]
df_adult = df_adult[which(df_adult$delivery!=777777 & df_adult$delivery!=999999),]

df_adult = df_adult[which(df_adult$alcoholfrq != 999),]

df_adult = df_adult[which(df_adult$meals_nothome != 5555 & df_adult$meals_nothome != 7777 & df_adult$meals_fastfood != 5555 & df_adult$meals_fastfood != 7777 & df_adult$meals_nothome != 7777 & df_adult$meals_fastfood != 7777),]

df_adult$tv_hrs[which(df_adult$tv_hrs == 0)] = 1
df_adult$tv_hrs[which(df_adult$tv_hrs == 8)] = 0
df_adult = df_adult[which(df_adult$tv_hrs != 77 & df_adult$tv_hrs != 99),]

df_adult = df_adult[which(df_adult$sleep_hr != 99),]

df_adult$smoke[which(is.na(df_adult$smoke))] = "missing"

# drop observations where number of missing value in certain columns is < 100
drop_obs = c("famincome", "grocery", "eatout", "delivery",
             "sleep_hr", "depressed", "sleep_trouble", "bmi")

for (feature in drop_obs){
  df_adult = df_adult[!is.na(df_adult[feature]),]
}

categorical_features = c("gender", "race", "edu", "marriage",
                        "famincome", "alcohol12yr", "milk", "depressed",
                        "sleep_trouble", "smoke", "tv_hrs")

numeric_features = c("age", "famsize", "alcoholfrq", "grocery", "eatout",
                    "delivery", "meals_nothome", "meals_fastfood", "sleep_hr")

# convert categorical variables into factors
df_adult[categorical_features] = lapply(df_adult[categorical_features], factor)

apply(df_adult, 2, function(x) sum(is.na(x))) # check how many missing data

```

```

##      gender      age      race      edu      marriage
##      0          0          0          0          0
##      famsize    famincome alcohol12yr alcoholfrq    grocery
##      0          0          0          0          0
##      eatout     delivery    milk    meals_nothome meals_fastfood
##      0          0          0          0          0
##      depressed  sleep_trouble tv_hrs    sleep_hr    smoke
##      0          0          0          0          0
##      bmi
##      0

```

```
sapply(df_adult, class) # check data classes
```

```

##      gender      age      race      edu      marriage
##      "factor"    "integer" "factor"    "factor"    "factor"
##      famsize    famincome alcohol12yr alcoholfrq    grocery
##      "integer"    "factor"    "factor"    "integer"    "integer"

```

```
##      eatout      delivery      milk  meals_nothome  meals_fastfood
##      "integer"    "integer"    "factor"    "integer"    "integer"
##      depressed  sleep_trouble    tv_hrs    sleep_hr      smoke
##      "factor"    "factor"    "factor"    "integer"    "factor"
##      bmi
##      "numeric"
```

Exploratory Data Analysis

Response Variable (BMI)

```
library(ggplot2)
require(gridExtra)

## Loading required package: gridExtra


##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##      combine

transform1 = log(df_adult$bmi)
transform2 = log(log(df_adult$bmi))

# response variable distribution
plot1 = ggplot(df_adult, aes(bmi)) +
  geom_histogram() +
  ggtitle("BMI")
plot2 = ggplot(df_adult, aes(transform1)) +
  geom_histogram() +
  ggtitle("Log(BMI)")
plot3 = ggplot(df_adult, aes(transform2)) +
  geom_histogram() +
  ggtitle("Log(Log(BMI))")
grid.arrange(plot1, plot2, plot3, ncol=3)

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

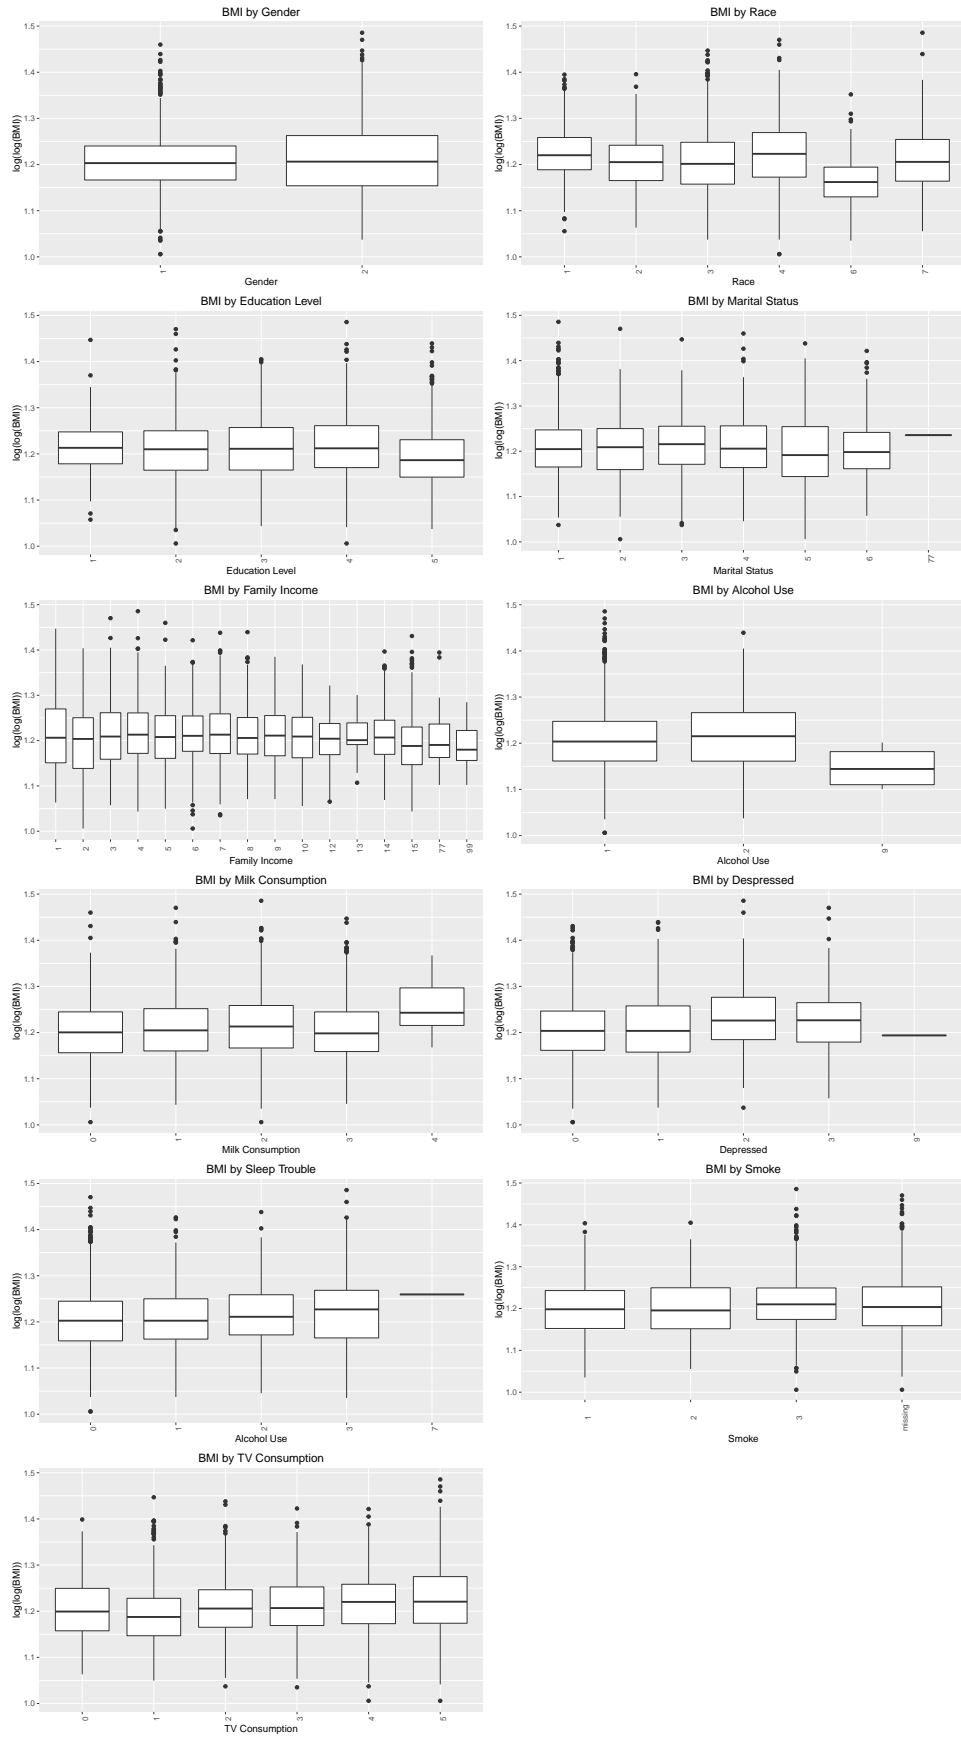


```
df_adult$log2bmi = log(log(df_adult$bmi))
```

Predictor Variables

Categorical Variables

```
plot1 = ggplot(df_adult, aes(x=gender, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Gender", y="log(log(BMI))") +  
  ggtitle("BMI by Gender")  
plot2 = ggplot(df_adult, aes(x=race, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Race", y="log(log(BMI))") +  
  ggtitle("BMI by Race")  
plot3 = ggplot(df_adult, aes(x=edu, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Education Level", y="log(log(BMI))") +  
  ggtitle("BMI by Education Level")  
plot4 = ggplot(df_adult, aes(x=marriage, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Marital Status", y="log(log(BMI))") +  
  ggtitle("BMI by Marital Status")  
plot5 = ggplot(df_adult, aes(x=famincome, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Family Income", y="log(log(BMI))") +  
  ggtitle("BMI by Family Income")  
plot6 = ggplot(df_adult, aes(x=alcohol12yr, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Alcohol Use", y="log(log(BMI))") +  
  ggtitle("BMI by Alcohol Use")  
plot7 = ggplot(df_adult, aes(x=milk, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Milk Consumption", y="log(log(BMI))") +  
  ggtitle("BMI by Milk Consumption")  
plot8 = ggplot(df_adult, aes(x=depressed, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Depressed", y="log(log(BMI))") +  
  ggtitle("BMI by Despressed")  
plot9 = ggplot(df_adult, aes(x=sleep_trouble, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Alcohol Use", y="log(log(BMI))") +  
  ggtitle("BMI by Sleep Trouble")  
plot10 = ggplot(df_adult, aes(x=smoke, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="Smoke", y="log(log(BMI))") +  
  ggtitle("BMI by Smoke")  
plot11 = ggplot(df_adult, aes(x=tv_hrs, y = log2bmi)) + geom_boxplot() +  
  theme(axis.text.x = element_text(angle=90)) +  
  labs(x="TV Consumption", y="log(log(BMI))") +  
  ggtitle("BMI by TV Consumption")  
grid.arrange(plot1, plot2, plot3, plot4, plot5, plot6, plot7, plot8, plot9, plot10, plot11, ncol=2)
```

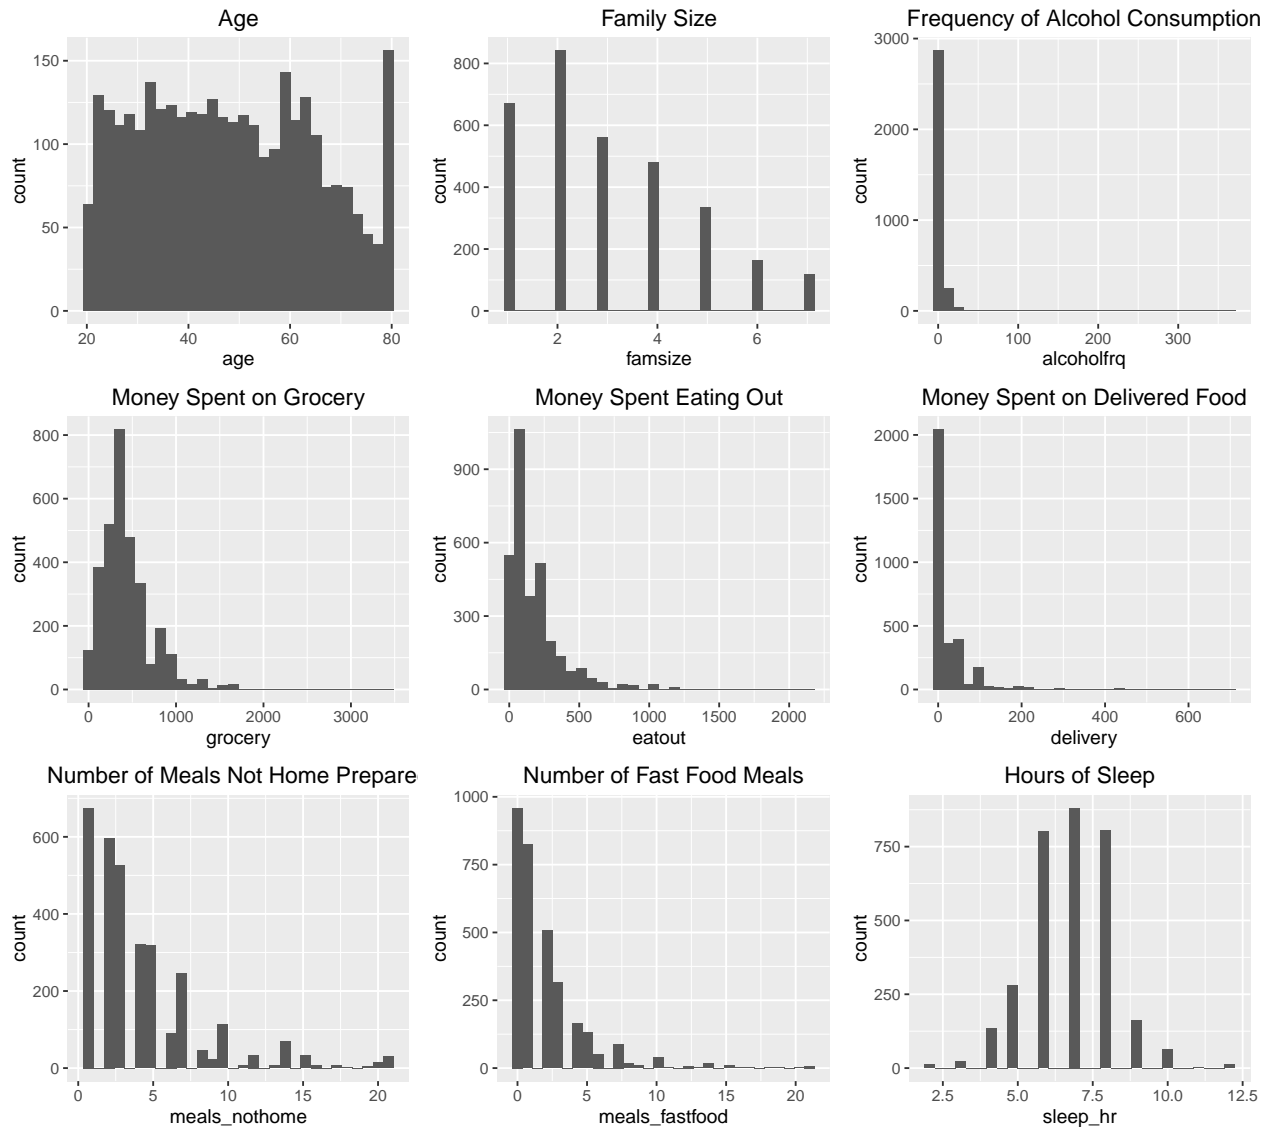


Numeric Variables

Distribtuion of numeric variables

```
# numeric predictor variable distributions
plot1 = ggplot(df_adult, aes(age)) + geom_histogram() +
  ggtitle("Age")
plot2 = ggplot(df_adult, aes(famsize)) + geom_histogram() +
  ggtitle("Family Size")
plot3 = ggplot(df_adult, aes(alcoholfrq)) + geom_histogram() +
  ggtitle("Frequency of Alcohol Consumption")
plot4 = ggplot(df_adult, aes(grocery)) + geom_histogram() +
  ggtitle("Money Spent on Grocery")
plot5 = ggplot(df_adult, aes(eatout)) + geom_histogram() +
  ggtitle("Money Spent Eating Out")
plot6 = ggplot(df_adult, aes(delivery)) + geom_histogram() +
  ggtitle("Money Spent on Delivered Food")
plot7 = ggplot(df_adult, aes(meals_nothome)) + geom_histogram() +
  ggtitle("Number of Meals Not Home Prepared")
plot8 = ggplot(df_adult, aes(meals_fastfood)) + geom_histogram() +
  ggtitle("Number of Fast Food Meals")
plot9 = ggplot(df_adult, aes(sleep_hr)) + geom_histogram() +
  ggtitle("Hours of Sleep")
grid.arrange(plot1, plot2, plot3, plot4, plot5, plot6, plot7, plot8, plot9, ncol=3)

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



Response vs. numeric distribution

```
# numeric predictor variable distributions
plot1 = ggplot(df_adult, aes(x=age, y=log2bmi)) + geom_point() +
  ggtitle("Age")
plot2 = ggplot(df_adult, aes(x=famsize, y=log2bmi)) + geom_point() +
  ggtitle("Family Size")
plot3 = ggplot(df_adult, aes(x=alcoholfrq, y=log2bmi)) + geom_point() +
  ggtitle("Frequency of Alcohol Consumption")
plot4 = ggplot(df_adult, aes(x=grocery, y=log2bmi)) + geom_point() +
  ggtitle("Money Spent on Grocery")
plot5 = ggplot(df_adult, aes(x=eatout, y=log2bmi)) + geom_point() +
  ggtitle("Money Spent Eating Out")
plot6 = ggplot(df_adult, aes(x=delivery, y=log2bmi)) + geom_point() +
  ggtitle("Money Spent on Delivered Food")
plot7 = ggplot(df_adult, aes(x=meals_nothome, y=log2bmi)) + geom_point() +
  ggtitle("Number of Meals Not Home Prepared")
plot8 = ggplot(df_adult, aes(x=meals_fastfood, y=log2bmi)) + geom_point() +
  ggtitle("Number of Fast Food Meals")
```

```
plot9 = ggplot(df_adult, aes(x=sleep_hr, y=log2bmi)) + geom_point() +
  ggtitle("Hours of Sleep")
grid.arrange(plot1, plot2, plot3, plot4, plot5, plot6, plot7, plot8, plot9, ncol=3)
```

